

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

AI Ethics & Safety (1 articles)

Samenvatting

OpenAI heeft laten zien hoe ze hun eigen AI-systemen in de gaten houden om te controleren of ze zich nog netjes gedragen. Stel je voor dat je een nieuwe medewerker hebt aangenomen: je kijkt niet alleen naar het eindresultaat van hun werk, maar je leest ook hun kladnotities om te begrijpen **hoe** ze tot hun beslissingen komen. Precies zo bestudeert OpenAI het "denkproces" van hun AI - de zogenaamde *chain-of-thought* - om vroegtijdig te ontdekken of een systeem ongewenst of onveilig gedrag vertoont. Het gaat specifiek om AI-programma's die zelfstandig code schrijven, waarbij het belangrijk is dat ze niet stiekem dingen doen die niemand heeft gevraagd. Dit soort transparantie en zelfcontrole wordt steeds belangrijker naarmate AI-systemen meer taken zelfstandig uitvoeren in de echte wereld.

How we monitor internal coding agents for misalignment

How OpenAI uses chain-of-thought monitoring to study misalignment in internal coding agents-analyzing real-world deployments to detect risks and strengthen AI safety safeguards.

OpenAI B

<https://openai.com>

AI Tools & Products (23 articles)

Samenvatting

OpenAI was deze week de grote speler: het bedrijf kondigde via zijn nieuwe stichting aan minstens een miljard dollar te investeren in gezondheidszorg en maatschappelijke projecten, en nam het bedrijf Astral over om betere hulpmiddelen voor softwareontwikkelaars te bouwen. Ondertussen geven techbedrijven hun AI-assistenten steeds meer zelfstandigheid - een beetje zoals een kind dat langzaam meer vrijheid krijgt, maar nog wél binnen afgesproken regels moet blijven, zoals Anthropic liet zien met zijn vernieuwde Claude Code. Opvallend was ook de aandacht voor veiligheid en controle: Spotify testte een hulpmiddel waarmee artiesten zich kunnen beschermen tegen nepnummers die door AI zijn gemaakt, en OpenAI bracht gratis bouwstenen uit waarmee app-makers hun producten veiliger kunnen maken voor tieners. Tot slot zagen we creatieve nieuwkomers zoals Talat en Hark, die AI dicht bij het dagelijks leven brengen met slimme vergadernotities en persoonlijke assistenten die jouw privacy vooropstellen.

Update on the OpenAI Foundation

The OpenAI Foundation announces plans to invest at least \$1 billion in curing diseases, economic opportunity, AI resilience, and community programs.

OpenAI B

<https://openai.com>

OpenAI to acquire Astral

Accelerates Codex growth to power the next generation of Python developer tools

OpenAI B

<https://openai.com>

What's New in Mellea 0.4.0 + Granite Libraries Release

<https://huggingface.co/blog/libm-granite/granite-libraries>

Hugging F

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

The hardest question to answer about AI-fueled delusions

MIT Techn

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. I was originally going to write this week's newsletter about AI and Iran, particularly the news we broke last Tuesday that the Pentagon is making plans for AI companies to train on

<https://www>

The Bay Area's animal welfare movement wants to recruit AI

MIT Techn

In early February, animal welfare advocates and AI researchers gathered in stocking feet at Mox, a scrappy, shoes-free coworking space in San Francisco. Yellow and red canopies billowed overhead, Persian rugs blanketed the floor, and mosaic lamps glowed beside potted plants. In the common area, a wildlife advocate spoke passionately to a crowd lounging in

<https://www>

Anthropic hands Claude Code more control, but keeps it on a leash

TechCrun

Anthropic's new auto mode for Claude Code lets AI execute tasks with fewer approvals, reflecting a broader shift toward more autonomous tools that balance speed with safety through built-in safeguards.

<https://tech>

Spotify tests new tool to stop AI slop from being attributed to real artists

TechCrun

The idea behind the new tool is to give artists more control over which tracks are associated with their name on Spotify.

<https://tech>

OpenAI adds open source tools to help developers build for teen safety

TechCrun

Rather than working from scratch to figure out how to make AI safer for teens, developers can use these policies to fortify what they build.

<https://tech>

Talat's AI meeting notes stay on your machine, not in the cloud

TechCrun

The subscription-free AI meeting notes app is a local-first twist on notetaking tools like Granola.

<https://tech>

Meet the former Apple designer building a new AI interface at Hark

TechCrun

The company said it would design models, hardware, and interfaces in tandem to deliver a "seamless end-to-end personal intelligence product."

<https://tech>

Air Street becomes one of the largest solo VCs in Europe with \$232M fund

TechCrun

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

London's Air Street Capital has raised a large Fund III with eyes locked on backing early-stage European and North American AI companies.

<https://tech>

Apple sets June date for WWDC 2026, teasing 'AI advancements'

TechCrunch

Apple will host its next Worldwide Developers Conference the week of June 8. The company is expected to announce major updates to Siri with advanced AI capabilities.

<https://tech>

How to Design a Production-Ready AI Agent That Automates Google Colab Workflows Using Colab-MCP, MCP Tools, FastMCP, and Kernel Execution

MarkTech

In this tutorial, we build an advanced, hands-on tutorial around Google's newly released colab-mcp, an open-source MCP (Model Context Protocol) server that lets any AI agent programmatically control Google Colab notebooks and runtimes. Across five self-contained snippets, we go from first principles to production-ready patterns. We start by constructing a minimal MCP tool registry from [82...

<https://www>

Decision-Making for Busy Owners: Moving Fast vs. Moving Slow (Sponsored)

KDNugget

Was "move fast and break things" ever really the best approach to building a business?

<https://bit.ly>

10 Best X (Twitter) Accounts to Follow for LLM Updates

KDNugget

Skip the hype and follow these 10 X accounts for reliable LLM papers, product launches, and thoughtful takes on where AI is heading.

<https://www>

5 Tips to Turn OpenAI Codex Into a Powerful AI Coding Agent

KDNugget

Learn five practical ways to make OpenAI Codex more agentic, reliable, and useful for real software engineering workflows.

<https://www>

5 Production Scaling Challenges for Agentic AI in 2026

ML Mastery

Everyone's <https://machinelearningmastery>.

<https://mac>

Top 46 AI Tools in 2026 You Must Use

Analytics

Now that we know AI is inevitably a part of our workflow, the more relevant question today is not "should I use AI?", but "how to use AI?". With the AI tools market more crowded than ever, each passing week sees a new assistant, generator, or automation. The struggle then is of choice from a [230] The post Top 46 AI Tools in 2026 You

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

Must Use appeared first on Analytics...

<https://www>

Top 5 Free Google Certificate Courses in 2026

Analytics

For different learning goals and career paths, choosing the right certification can get confusing. Some people want analytics. Others want ads. Some care about AI. And many just want something credible to add to their resume. This list is built with that in mind. A set of free Google certificate courses, each aligned to a […] The post Top 5 Free Google Certificate Courses in 2026 appeared fi...

<https://www>

Guide to Propensity Score Matching for Causal Inference to Estimate True Impact

Analytics

One of the core challenges of data science is drawing meaningful causal conclusions from observational data. In many such cases, the goal is to estimate the true impact of a treatment or behaviour as fairly as possible. This article explores Propensity Score Matching (PSM), a statistical technique used for that very purpose. Unlike randomized experiments […] The post Guide to Propensity Scor...

<https://www>

Top 10 AI Coding Assistants of 2026

Analytics

AI coding assistants have quickly moved from optional tools to a core part of modern software development. Adoption is accelerating fast. Around 84% of developers now use or plan to use AI tools, and over half use them daily. The market has already reached about \$8.5 billion in 2026 and is growing rapidly. These tools are not just helping […] The post Top 10 AI Coding Ass...

<https://www>

Inside Claude Cowork: How to Run Agentic AI Tasks Like a Pro

Analytics

Most AI tools still require constant supervision, forcing you to guide every step. Claude Cowork,, the latest offering by Anthropic, changes that! By bringing an agentic system into everyday workflows, you describe the outcome and let it handle the execution independently. It can deliver organized files, structured documents, and synthesized research while you focus elsewhere, […] The post I...

<https://www>

Top 7 Free Data Analytics Courses with Certificates

Analytics

For different learning styles, career goals, and comfort with tools, finding the right data analyst course is HARD. Some people start with Excel. Others into Python. With no clear roadmap ahead, it's hard to find a single starting point. Some want job-ready certifications, while others just want hands-on practice. This list is built for that. […] The post Top 7 Free Data Analytics Cour...

<https://www>

AI in Business (7 articles)

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

Samenvatting

Grote investeerders gooien miljarden in bedrijven die met kunstmatige intelligentie werken - alsof ze goudkoorts hebben en AI de nieuwe goudader is. Kleiner Perkins haalde 3,5 miljard dollar op om in AI-bedrijven te investeren, en Databricks gebruikt zijn volle spaarpot om slimme startups op te kopen die hun producten veiliger maken. Ondertussen weigerde een boerin uit Kentucky 26 miljoen dollar van een groot AI-bedrijf dat op haar land een enorm computergebouw wilde neerzetten. Aan de praktische kant zien we dat AI steeds meer alledaagse bedrijfstaken overneemt, van voorraadbeheer in magazijnen tot het bouwen van software en het aansturen van robots. De boodschap is duidelijk: AI dringt razendsnel door in het bedrijfsleven, en bedrijven investeren fors om daar een plekje vooraan te bemachtigen.

With \$3.5B in fresh capital, Kleiner Perkins is going all in on AI

TechCrunch

The fundraise includes \$1 billion for investing in early-stage startups, and \$2.5 billion for late-stage growth businesses.

<https://techcrunch.com/2026/03/25/kleiner-perkins-raises-35-billion-for-ai-investments/>

Kentucky woman rejects \$26M offer to turn her farm into a data center

TechCrunch

A "major artificial intelligence company" reportedly offered a Kentucky family \$26 million to build a data center on their farm.

<https://techcrunch.com/2026/03/25/kentucky-woman-rejects-26-million-offer-to-turn-farm-into-data-center/>

Databricks bought two startups to underpin its new AI security product

TechCrunch

With an overflowing war chest from its recent \$5 billion raise, Databricks is buying startups and looking for more. It acquired Antimatter and SiftD.ai.

<https://techcrunch.com/2026/03/25/databricks-bought-two-startups-to-underpin-its-new-ai-security-product/>

Doss raises \$55M for AI inventory management that plugs into ERP

TechCrunch

Doss's AI-powered inventory management system integrates with existing ERP systems. The Series B round was co-led by Madrona and Premji Invest.

<https://techcrunch.com/2026/03/25/doss-raises-55-million-for-ai-inventory-management-that-plugs-into-erp/>

Agile Robots becomes the latest robotics company to partner with Google DeepMind

TechCrunch

Agile Robots will incorporate Google DeepMind's robotics foundation models into its bots while collecting data for the AI research lab.

<https://techcrunch.com/2026/03/25/agile-robots-becomes-the-latest-robotics-company-to-partner-with-google-deepmind/>

Vibe-coding startup Lovable is on the hunt for acquisitions

TechCrunch

Lovable's founder said the fast-growing vibe-coding startup is looking for startups and teams to join its company.

<https://techcrunch.com/2026/03/25/vibe-coding-startup-lovable-is-on-the-hunt-for-acquisitions/>

Analytics Patterns Every Data Scientist Should Master

KDnuggets

Learn the analytics pattern you can use in most business analytics tasks.

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

<https://www>

Image & Video AI (5 articles)

Samenvatting

OpenAI heeft zijn videogenerator Sora - waarmee je met een simpele tekstopdracht levensechte video's kon maken - alweer stopgezet, omdat te weinig mensen het bleven gebruiken. Het was als een spectaculair nieuw pretpark dat na de openingsweek al leegliep. Tegelijkertijd lanceerde Luma Labs een nieuw AI-model genaamd Uni-1, dat eerst "nadenkt" over wat je bedoelt voordat het een afbeelding maakt, in plaats van blindelings pixels samen te plakken. Google zet ondertussen in op veiligheid met SynthID, een technologie die onzichtbare watermerken in AI-gemaakte beelden, tekst, audio en video verstopt, zodat je altijd kunt controleren of iets door een mens of machine is gemaakt. De rode draad: de wereld van AI-gegenereerde beelden en video's ontwikkelt zich razendsnel, maar het vinden van een publiek én het veilig houden ervan blijken minstens zo'n grote uitdaging als de technologie zelf.

Creating with Sora Safely

OpenAI B

To address the novel safety challenges posed by a state-of-the-art video model as well as a new social creation platform, we've built Sora 2 and the Sora app with safety at the foundation. Our approach is anchored in concrete protections.

<https://open>

OpenAI's Sora was the creepiest app on your phone - now it's shutting down

TechCrun

Though the underlying Sora 2 video- and audio-generation model is scarily impressive, there was not sustained interest in an AI-only social feed.

<https://tech>

Luma Labs Launches Uni-1: The Autoregressive Transformer Model that Reasons through Intentions Before Generating Images

MarkTech

In the field of generative AI media, the industry is transitioning from purely probabilistic pixel synthesis toward models capable of structural reasoning. Luma Labs has just released Uni-1, a foundational image model designed to address the [intent gap](#); inherent in standard diffusion pipelines. By implementing a reasoning phase prior to generation, Uni-1 shifts the workflow [\[8230;\]](#) ...

<https://www>

SynthID: What it is and How it Works

KDnugget

Learn everything about SynthID, how it embeds invisible AI watermarks, and how it verifies and identifies AI-generated content across text, images, audio, and video

<https://www>

Top 10 YouTube Channels to Learn Machine Learning

Analytics

With so much happening in AI and machine learning today, figuring out where to start can feel overwhelming. Different learners prefer different approaches! Some want visuals, others prefer coding. Some prefer short form, others lean toward

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

long-form learning. While many simply want a clear path into ML. This article is here to fix that. Instead [8230:] The post Top 10 YouTube Channels to Learn M...

<https://www>

Large Language Models (20 articles)

Samenvatting

OpenAI probeerde van ChatGPT een soort digitaal winkelcentrum te maken waar je direct producten kon kopen, maar dat plan werd al snel weer teruggedraaid - winkelen via een chatbot bleek toch lastiger dan gedacht. Ondertussen waaide er een stevige politieke storm rond AI: het Amerikaanse leger wil kunstmatige intelligentie inzetten voor defensie, wat leidde tot ruzie met Anthropic en een omstreden deal met OpenAI, terwijl in Londen duizenden mensen de straat op gingen om tegen AI te protesteren. Op technisch vlak werkten onderzoekers aan manieren om grote taalmodellen - de motoren achter chatbots zoals ChatGPT - zuiniger en slimmer te laten draaien, vergelijkbaar met een auto die leert om dezelfde afstand met minder brandstof af te leggen. Google bracht zijn eigen AI-assistent Gemini naar televisies, en OpenAI zette stappen om chatbots veiliger te maken voor tieners. Kortom: de AI-wereld groeit razendsnel, maar botst steeds vaker op praktische, ethische en maatschappelijke grenzen.

Helping developers build safer AI experiences for teens

OpenAI B

OpenAI releases prompt-based teen safety policies for developers using gpt-oss-safeguard, helping moderate age-specific risks in AI systems.

<https://open>

Powering product discovery in ChatGPT

OpenAI B

ChatGPT introduces richer, visually immersive shopping powered by the Agentic Commerce Protocol, enabling product discovery, side-by-side comparisons, and merchant integration.

<https://open>

The AI Hype Index: AI goes to war

MIT Techn

AI is at war. Anthropic and the Pentagon feuded over how to weaponize Anthropic's AI model Claude; then OpenAI swept the Pentagon off its feet with an "opportunistic and sloppy" deal. Users quit ChatGPT in droves. People marched through London in the biggest protest against AI to date. If you're keeping score, Anthropic-the company founded8230;

<https://www>

OpenAI's plans to make ChatGPT more like Amazon aren't going so well

TechCrun

OpenAI says it's moving away from Instant Checkout, which allowed users to buy items directly through the ChatGPT interface.

<https://tech>

Google TV's new Gemini features keep fans updated on sports teams and more

TechCrun

Three Gemini-powered features are coming to your Google TV. This includes visual responses, deep dives, and sports briefs.

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

<https://tech>

Mirage raises \$75M to continue building models for its AI video-editing app Captions

TechCrunch

Mirage, the maker of video-editing app Captions, has raised \$75 million in growth financing from General Catalyst's Customer Value Fund (CVF).

<https://tech>

Bernie Sanders' AI 'gotcha' video flops, but the memes are great

TechCrunch

Sen. Bernie Sanders thinks he's tricked Claude into revealing the AI industry's secrets, but he really just exposed how agreeable chatbots can become.

<https://tech>

NVIDIA AI Introduces PivotRL: A New AI Framework Achieving High Agentic Accuracy With 4x Fewer Rollout Turns Efficiently

MarkTechInsights

Post-training Large Language Models (LLMs) for long-horizon agentic tasks-such as software engineering, web browsing, and complex tool use-presents a persistent trade-off between computational efficiency and model generalization. While Supervised Fine-Tuning (SFT) is computationally inexpensive, it frequently suffers from out-of-domain (OOD) performance degradation and struggles to generalize beyo...

<https://www>

Paged Attention in Large Language Models LLMs

MarkTechInsights

When running LLMs at scale, the real limitation is GPU memory rather than compute, mainly because each request requires a KV cache to store token-level data. In traditional setups, a large fixed memory block is reserved per request based on the maximum sequence length, which leads to significant unused space and limits concurrency. Paged Attention [1] The post Paged Attention in Large Language...

<https://www>

A Coding Implementation to Design Self-Evolving Skill Engine with OpenSpace for Skill Learning, Token Efficiency, and Collective Intelligence

MarkTechInsights

In this tutorial, we explore OpenSpace, a self-evolving skill engine developed by HKUDS that makes AI agents smarter, more cost-efficient, and capable of learning from every task they perform. We walk through the complete lifecycle of OpenSpace: from installing and configuring an OpenAI model, to executing cold-start tasks where no prior skills exist, watching the [1] The post A Coding Implementation...

<https://www>

This AI Paper Introduces TinyLoRA, A 13-Parameter Fine-Tuning Method That Reaches 91.8 Percent GSM8K on Qwen2.5-7B

MarkTechInsights

Researchers from FAIR at Meta, Cornell University, and Carnegie Mellon University have demonstrated that large language models (LLMs) can learn to reason using a remarkably small number of trained parameters. The research team introduces

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

TinyLoRA, a parameterization that can scale down to a single trainable parameter under extreme sharing settings. Using this method on a [8230;] The post This AI...

<https://www>

Meta AI's New Hyperagents Don't Just Solve Tasks-They Rewrite the Rules of How They Learn

MarkTech

The dream of recursive self-improvement in AI-where a system doesn't just get better at a task, but gets better at learning-has long been the 'holy grail' of the field. While theoretical models like the Gödel Machine have existed for decades, they remained largely impractical in real-world settings. That changed with the Darwin Gödel Machine (DGM), [8230;] The post Meta AI's New Hype...

<https://www>

How BM25 and RAG Retrieve Information Differently?

MarkTech

When you type a query into a search engine, something has to decide which documents are actually relevant - and how to rank them. BM25 (Best Matching 25), the algorithm powering search engines like Elasticsearch and Lucene, has been the dominant answer to that question for decades. It scores documents by looking at three things: [8230;] The post How BM25 and RAG Retrieve Information Differ...

<https://www>

ChatLLM Review: Tired of Multiple AI Tools? Here's a Smarter All-in-One Alternative

KDnugget

Explore ChatLLM by Abacus AI, an all-in-one AI platform that brings together tools like ChatGPT, Claude, and Midjourney into a single workflow. Learn about its features, pricing, and real-world use cases.

<https://www>

Getting Started with Nanobot: Build Your First AI Agent

KDnugget

Learn how to set up Nanobot, connect it to WhatsApp, and power it with OpenAI GPT-5.3-Codex for a practical, always-on AI agent.

<https://www>

Beyond the Vector Store: Building the Full Data Layer for AI Applications

ML Master

If you look at the architecture diagram of almost any AI startup today, you will see a large language model (LLM) connected to a vector store.

<https://mac>

7 Readability Features for Your Next Machine Learning Model

ML Master

Unlike fully structured tabular data, preparing text data for machine learning models typically entails tasks like tokenization, embeddings, or sentiment analysis.

<https://mac>

Mistral Small 4: The One Model That Codes, Reasons, and Chats

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

Analytics

Artificial intelligence is rapidly evolving. New models emerge nearly every day, with each one attempting to be the best. In this sea of similar models, we see something new every now and then. One of such models is the new Mistral Small 4. It is an innovative AI model that is not only going to […] The post Mistral Small 4: The One Model That Codes, Reasons, and Chats appeared first on Analy...

<https://www>

PageIndex vs Traditional RAG: A Better Way to Build Document Chatbots

Analytics

What if the way we build AI document chatbots today is flawed? Most systems use RAG. They split documents into chunks, create embeddings, and retrieve answers using similarity search. It works in demos but often fails in real use. It misses obvious answers or picks the wrong context. Now there is a new approach called […] The post PageIndex vs Traditional RAG: A Better Way to Build Document ...

<https://www>

Claude Skills Explained: Build, Configure, and Use Custom Skills on Claude Code

Analytics

If you regularly use AI, especially for coding, you know there is an obvious upgrade to the usual to-and-fro in chats. What if you could save a particular workflow in AI and run it without writing a super-long prompt every time you need it? Claude Skills now lets you do exactly that. This nifty little […] The post Claude Skills Explained: Build, Configure, and Use Custom Skills on Claude Cod...

<https://www>

Research & Science (3 articles)

Samenvatting

OpenAI zet al zijn kaarten op één paard: het bedrijf wil een volledig zelfstandige AI-onderzoeker bouwen - een soort digitale wetenschapper die zonder menselijke hulp grote, ingewikkelde problemen kan oplossen. Google heeft ondertussen een slimme truc bedacht om AI-modellen zuiniger met geheugen te laten omgaan, waardoor ze tot zes keer minder opslagruimte nodig hebben zonder dat de kwaliteit achteruitgaat - vergelijk het met een koffer die je zes keer kleiner kunt vouwen terwijl er nog steeds al je kleding in past. En AI-onderzoeker Yann LeCun werkt aan een nieuw type model dat de wereld leert begrijpen door naar beelden te kijken, maar worstelt met het probleem dat zo'n model soms lui wordt en steeds dezelfde voorspelling herhaalt in plaats van echt te leren. Samen laten deze ontwikkelingen zien dat de AI-wereld op volle kracht vooruit denderd, maar ook tegen serieuze technische grenzen aanloopt.

OpenAI is throwing everything into building a fully automated researcher

MIT Techn

OpenAI is refocusing its research efforts and throwing its resources into a new grand challenge. The San Francisco firm has set its sights on building what it calls an AI researcher, a fully automated agent-based system that will be able to go off and tackle large, complex problems by itself. ??OpenAI says that this new…

<https://www>

Google Introduces TurboQuant: A New Compression Algorithm that Reduces LLM Key-Value Cache Memory by 6x and Delivers Up to 8x Speedup, All with Zero Accuracy Loss

MarkTech

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

The scaling of Large Language Models (LLMs) is increasingly constrained by memory communication overhead between High-Bandwidth Memory (HBM) and SRAM. Specifically, the Key-Value (KV) cache size scales with both model dimensions and context length, creating a significant bottleneck for long-context inference. Google research team has proposed TurboQuant, a data-oblivious quantization framework des...

<https://www>

Yann LeCun's New LeWorldModel (LeWM) Research Targets JEPAs Collapse in Pixel-Based Predictive World Modeling

MarkTech

World Models (WMs) are a central framework for developing agents that reason and plan in a compact latent space. However, training these models directly from pixel data often leads to "representation collapse," where the model produces redundant embeddings to trivially satisfy prediction objectives. Current approaches attempt to prevent this by relying on complex heuristics: they

<https://www>

Robotics & Hardware (2 articles)

Samenvatting

Drones die ramen wassen klinkt als sciencefiction, maar het bedrijf Lucid Bots doet het gewoon - en de vraag is zo groot dat ze 20 miljoen dollar ophaalden om bij te kunnen blijven. Denk aan een vliegende glazenwasser die zonder hoogtevrees de hoogste wolkenkrabbers te lijf gaat. Aan de andere kant maakte chipontwerper Arm ook groot nieuws: na 35 jaar tekeningen maken voor anderen, bouwt het bedrijf nu voor het eerst zelf een chip - een beetje alsof een architect eindelijk zijn eigen droomhuis bouwt in plaats van alleen blauwdrukken te leveren. De eerste klant is techgigant Meta, het moederbedrijf van Facebook. Kortom, de robotica- en hardwarewereld zet flinke stappen vooruit, zowel in de lucht als in de fabriek.

Lucid Bots raises \$20M to keep up with demand for its window-washing drones

TechCrunch

Lucid Bots has seen demand accelerate over the last year for its window cleaning drones and power washing robots.

<https://tech>

Arm is releasing the first in-house chip in its 35-year history

TechCrunch

Arm is producing its own CPU for the first time. It developed the CPU with Meta, which is also the chip's first customer.

<https://tech>

General AI News (8 articles)

Samenvatting

AI-systemen worden steeds zelfstandiger: in plaats van alleen antwoorden te geven, kunnen ze nu ook daadwerkelijk taken uitvoeren, zoals een complete vakantie boeken en afrekenen - denk aan een digitale persoonlijke assistent die niet alleen advies geeft, maar ook de telefoon pakt en alles voor je regelt. Tegelijkertijd werken ontwikkelaars hard aan de "hersenen" achter deze slimme systemen, met nieuwe manieren om te meten hoe goed AI-stemassistenten presteren en hoe je ze een beter geheugen kunt geven. Maar net als een auto die soms onverwacht een verkeerde afslag neemt, blijken AI-agents ook te falen door kleine technische instellingen die hun gedrag onvoorspelbaar maken. Daarnaast verschijnen er veel praktische

Weekly AI News Digest

Generated on 2026-03-25 - Articles from the last 7 days

handleidingen om de programmeertaal Python - de belangrijkste taal achter AI - sneller en betrouwbaarder te maken. Kortom: de AI-wereld bouwt volop aan systemen die meer kunnen doen, maar worstelt tegelijk met het temmen van hun onvoorspelbare kanten.

A New Framework for Evaluating Voice Agents (EVA)

<https://huggingface.co/blog/ServiceNow-AI/eva>

Hugging F

Build a Domain-Specific Embedding Model in Under a Day

<https://huggingface.co/blog/nvidia/domain-specific-embedding-finetune>

Hugging F

Agentic commerce runs on truth and context

Imagine telling a digital agent, "Use my points and book a family trip to Italy. Keep it within budget, pick hotels we've liked before, and handle the details." Instead of returning a list of links, the agent assembles an itinerary and executes the purchase. That shift, from assistance to execution, is what makes agentic AI

MIT Techn

<https://www>

5 Useful DIY Python Functions for Error Handling

Debugging Python doesn't need to be complicated. These 5 DIY functions simplify error handling and improve code reliability.

KDnugget

<https://www>

How to Speed Up Slow Python Code Even If You're a Beginner

Slow Python code is often caused by small inefficiencies. With a few beginner-friendly techniques, you can make your programs run noticeably faster. This guide will teach you how.

KDnugget

<https://www>

5 Powerful Python Decorators for Robust AI Agents

These five Python decorators have saved me from countless headaches, and they will probably save you, too.

KDnugget

<https://www>

7 Steps to Mastering Memory in Agentic AI Systems

Memory is one of the most overlooked parts of agentic system design.

ML Maste

<https://mac>

Why Agents Fail: The Role of Seed Values and Temperature in Agentic Loops

In the modern AI landscape, an agent loop is a cyclic, repeatable, and continuous process whereby an entity called an AI agent — with a certain degree of autonomy — works toward a goal.

ML Maste

<https://mac>