

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

## AI Ethics & Safety (2 articles)

### Samenvatting

OpenAI, het bedrijf achter ChatGPT, vraagt nu buitenstaanders om mee te zoeken naar zwakke plekken in hun AI-systemen - een beetje zoals een autofabrikant die onafhankelijke experts uitnodigt om te controleren of de remmen wel goed werken. Ze belonen mensen die gevaarlijke fouten ontdekken, zoals manieren waarop kwaadwillenden de AI kunnen misbruiken of gevoelige informatie kunnen stelen. Ondertussen trekken de Amerikaanse politici Bernie Sanders en Alexandria Ocasio-Cortez aan de noodrem: zij willen dat de bouw van nieuwe datacenters - de enorme computergebouwen die AI draaiende houden - wordt stilgelegd totdat er duidelijke regels komen. Zij vinden dat de technologie zich sneller ontwikkelt dan de wetgeving kan bijhouden. Kortom: er wordt zowel vanuit de techsector als vanuit de politiek hard gewerkt om AI veiliger en beter gereguleerd te krijgen.

### Introducing the OpenAI Safety Bug Bounty program

OpenAI launches a Safety Bug Bounty program to identify AI abuse and safety risks, including agentic vulnerabilities, prompt injection, and data exfiltration.

OpenAI B

<https://openai.com/safety-bug-bounty>

### Bernie Sanders and AOC propose a ban on data center construction

Senator Bernie Sanders and Rep. Alexandria Ocasio-Cortez introduced companion legislation to halt construction on new data centers until Congress passes comprehensive AI regulation.

TechCrunch

<https://techcrunch.com/2026/03/26/bernie-sanders-aoc-propose-ban-on-data-center-construction/>

## AI Tools & Products (25 articles)

### Samenvatting

OpenAI publiceerde een soort "gedragshandleiding" voor hun AI-modellen - vergelijk het met een opvoedingsboek dat bepaalt hoe de AI zich hoort te gedragen, met regels rond veiligheid en vrijheid. Daarnaast kondigde de OpenAI Foundation aan minstens 1 miljard dollar te investeren in onder meer ziektebestrijding en economische kansen. Google bracht Lyria 3 Pro uit, een vernieuwde AI die langere en betere muziekstukken kan maken, en rolde die technologie uit over meerdere producten. Meta zet AI in om winkelen op Instagram en Facebook makkelijker te maken, bijvoorbeeld door je automatisch meer productinformatie te tonen. Ondertussen duiken er steeds creatievere toepassingen op: van een gratis wiskundetool die patronen opspoorde tot onderzoekers die AI willen inzetten voor dierenwelzijn.

### Inside our approach to the Model Spec

Learn how OpenAI's Model Spec serves as a public framework for model behavior, balancing safety, user freedom, and accountability as AI systems advance.

OpenAI Blog

<https://openai.com/model-spec>

### Update on the OpenAI Foundation

The OpenAI Foundation announces plans to invest at least \$1 billion in curing diseases, economic opportunity, AI resilience, and community programs.

OpenAI Blog

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

<https://openai.com>

## Lyria 3 Pro: Create longer tracks in more

Google De

Introducing Lyria 3 Pro, which unlocks longer tracks with structural awareness. We're also bringing Lyria to more Google products and surfaces.

<https://deepmind.com>

## What's New in Mellea 0.4.0 + Granite Libraries Release

Hugging F

<https://huggingface.co/blog/ibm-granite/granite-libraries>

## This startup wants to change how mathematicians do math

MIT Techn

Axiom Math, a startup based in Palo Alto, California, has released a free new AI tool for mathematicians, designed to discover mathematical patterns that could unlock solutions to long-standing problems. The tool, called Explorer, is a redesign of an existing one called PatternBoost that François Charton, now a research scientist at Axiom, co-developed in 2024

<https://www.aaxiommath.com>

## The hardest question to answer about AI-fueled delusions

MIT Techn

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. I was originally going to write this week's newsletter about AI and Iran, particularly the news we broke last Tuesday that the Pentagon is making plans for AI companies to train on

<https://www.thealgorithm.ai>

## The Bay Area's animal welfare movement wants to recruit AI

MIT Techn

In early February, animal welfare advocates and AI researchers gathered in stocking feet at Mox, a scrappy, shoes-free coworking space in San Francisco. Yellow and red canopies billowed overhead, Persian rugs blanketed the floor, and mosaic lamps glowed beside potted plants. In the common area, a wildlife advocate spoke passionately to a crowd lounging in

<https://www.moxsf.com>

## The least surprising chapter of the Manus story is what's happening right now

TechCrun

Did anyone think there would not be a reckoning over this tie-up?

<https://techcrunch.com>

## Google launches Lyria 3 Pro music generation model

TechCrun

Google is launching Lyria 3 Pro, an upgraded music model that generates longer, more customizable tracks, as it expands AI music tools across Gemini, enterprise products, and other services.

<https://techcrunch.com>

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

## Meta turns to AI to make shopping easier on Instagram and Facebook

Meta is using generative AI to provide more product and brand information to consumers when they're shopping in its apps.

TechCrunch

<https://techcrunch.com/2026/03/26/meta-ai-shopping/>

## With Sift, two ex-SpaceX engineers are bringing the software that helped launch rockets to the factory floor

Sift is building the data infrastructure for advanced manufacturing.

TechCrunch

<https://techcrunch.com/2026/03/26/sift-manufacturing/>

## Anthropic hands Claude Code more control, but keeps it on a leash

Anthropic's new auto mode for Claude Code lets AI execute tasks with fewer approvals, reflecting a broader shift toward more autonomous tools that balance speed with safety through built-in safeguards.

TechCrunch

<https://techcrunch.com/2026/03/26/anthropic-claude-code/>

## Spotify tests new tool to stop AI slop from being attributed to real artists

The idea behind the new tool is to give artists more control over which tracks are associated with their name on Spotify.

TechCrunch

<https://techcrunch.com/2026/03/26/spotify-ai-slop/>

## Building Declarative Data Pipelines with Snowflake Dynamic Tables: A Workshop Deep Dive

Traditional data pipeline development often requires extensive procedural code to define how data should be transformed and moved between stages. The declarative approach flips this paradigm by allowing data engineers to specify what the end result should be rather than prescribing every step of how to achieve it.

KDNugget

<https://www.kdnugget.com/2026/03/26/building-declarative-data-pipelines-with-snowflake-dynamic-tables-a-workshop-deep-dive/>

## Decision-Making for Busy Owners: Moving Fast vs. Moving Slow (Sponsored)

Was "move fast and break things" ever really the best approach to building a business?

KDNugget

<https://bit.ly/3QvXvXv>

## 10 Best X (Twitter) Accounts to Follow for LLM Updates

Skip the hype and follow these 10 X accounts for reliable LLM papers, product launches, and thoughtful takes on where AI is heading.

KDNugget

<https://www.kdnugget.com/2026/03/26/10-best-x-accounts-to-follow-for-llm-updates/>

## 5 Tips to Turn OpenAI Codex Into a Powerful AI Coding Agent

Learn five practical ways to make OpenAI Codex more agentic, reliable, and useful for real software engineering workflows.

KDNugget

<https://www.kdnugget.com/2026/03/26/5-tips-to-turn-openai-codex-into-a-powerful-ai-coding-agent/>

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

## 5 Production Scaling Challenges for Agentic AI in 2026

Everyone's <https://machinelearningmastery>.

ML Master

<https://mac>

## 20+ Solved AI Projects to Build Your Portfolio and Boost Your Resume

Analytics

Projects are the bridge between learning and becoming a professional. While theory builds fundamentals, recruiters value candidates who can solve real problems. A strong, diverse portfolio showcases practical skills, technical range, and problem-solving ability. This guide compiles over 20 solved projects across AI domains, from basic machine learning to advanced generative AI and agentic systems....

<https://www>

## Claude Dispatch: How to Use It for Remote AI Automation

Analytics

My main complaint with AI solutions is that they are largely dependent on my presence for any task. Even with agentic AI now in the mix, complete automation of any complex process still seems like a myth. Tools like n8n and make.com need a considerable setup time and do not really function as conventional AI [The post Claude Dispatch: How to Use It for Remote AI Automation appeared first ...

<https://www>

## Top 46 AI Tools in 2026 You Must Use

Analytics

Now that we know AI is inevitably a part of our workflow, the more relevant question today is not "should I use AI?", but "how to use AI?". With the AI tools market more crowded than ever, each passing week sees a new assistant, generator, or automation. The struggle then is of choice from a [The post Top 46 AI Tools in 2026 You Must Use appeared first on Analytics...

<https://www>

## Top 5 Free Google Certificate Courses in 2026

Analytics

For different learning goals and career paths, choosing the right certification can get confusing. Some people want analytics. Others want ads. Some care about AI. And many just want something credible to add to their resume. This list is built with that in mind. A set of free Google certificate courses, each aligned to a [The post Top 5 Free Google Certificate Courses in 2026 appeared fi...

<https://www>

## Guide to Propensity Score Matching for Causal Inference to Estimate True Impact

Analytics

One of the core challenges of data science is drawing meaningful causal conclusions from observational data. In many such cases, the goal is to estimate the true impact of a treatment or behaviour as fairly as possible. This article explores Propensity Score Matching (PSM), a statistical technique used for that very purpose. Unlike randomized experiments [The post Guide to Propensity Scor...

<https://www>

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

## Top 10 AI Coding Assistants of 2026

Analytics

AI coding assistants have quickly moved from optional tools to a core part of modern software development. Adoption is accelerating fast. Around 84% of developers now use or plan to use AI tools, and over half use them daily. The market has already reached about \$8.5 billion in 2026 and is growing rapidly. These tools are not just helping

[The post Top 10 AI Coding Ass...

<https://www>

## Inside Claude Cowork: How to Run Agentic AI Tasks Like a Pro

Analytics

Most AI tools still require constant supervision, forcing you to guide every step. Claude Cowork,, the latest offering by Anthropic, changes that! By bringing an agentic system into everyday workflows, you describe the outcome and let it handle the execution independently. It can deliver organized files, structured documents, and synthesized research while you focus elsewhere,

[The post I...

<https://www>

## AI in Business (8 articles)

### Samenvatting

Grote investeerders pompen miljarden in bedrijven die kunstmatige intelligentie bouwen - alsof ze wedden op de winnende paarden in een race die net begonnen is. Advocatenkantoor-app Harvey is nu 11 miljard dollar waard, vergader-app Granola groeide in één klap naar 1,5 miljard, en investeringsfonds Kleiner Perkins stak 3,5 miljard in AI-startups. Ondertussen waarschuwt Anthropic dat er een kloof groeit tussen mensen die goed met AI overweg kunnen en mensen die dat niet kunnen, wat op termijn voor ongelijkheid op de werkvloer kan zorgen. Meta wil kleine ondernemers helpen om AI te gebruiken, en Databricks kocht twee kleine bedrijfjes op om zijn klanten beter te beschermen tegen digitale dreigingen. En te midden van al dat tech-geweld weigerde een boerin uit Kentucky 26 miljoen dollar van een AI-bedrijf dat op haar land een datacenter wilde bouwen - niet alles is te koop.

## The AI skills gap is here, says AI company, and power users are pulling ahead

TechCrunch

Anthropic finds AI isn't replacing jobs yet, but early data shows growing inequality as experienced users gain an edge, raising concerns about future displacement and workforce divides.

<https://tech>

## Harvey confirms \$11B valuation: Sequoia triples down

TechCrunch

Investors like Sequoia, Andreessen Horowitz, Kleiner Perkins, and Elad Gil can't get enough of AI legal tech startup Harvey.

<https://tech>

## Granola raises \$125M, hits \$1.5B valuation as it expands from meeting notetaker to enterprise AI app

TechCrunch

Granola's valuation jumped from \$250 million to \$1.5 billion with this round, and it has added more support for AI agents after users previously complained.

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

<https://tech>

## Meta launches new initiative to support entrepreneurship, drive AI adoption

TechCrunch

Meta CEO Mark Zuckerberg said in a memo to staff that small businesses have always been a big part of the company's business model, and that while tens of millions of entrepreneurs already use its platforms to grow and connect with customers, the company wants to do more in the space.

<https://tech>

## With \$3.5B in fresh capital, Kleiner Perkins is going all in on AI

TechCrunch

The fundraise includes \$1 billion for investing in early-stage startups and \$2.5 billion for late-stage growth businesses.

<https://tech>

## Kentucky woman rejects \$26M offer to turn her farm into a data center

TechCrunch

A "major artificial intelligence company" reportedly offered a Kentucky family \$26 million to build a data center on their farm.

<https://tech>

## Databricks bought two startups to underpin its new AI security product

TechCrunch

With an overflowing war chest from its recent \$5 billion raise, Databricks is buying startups and looking for more. It acquired Antimatter and SiftD.ai.

<https://tech>

## Analytics Patterns Every Data Scientist Should Master

KDNuggets

Learn the analytics pattern you can use in most business analytics tasks.

<https://www>

## Image & Video AI (5 articles)

### Samenvatting

OpenAI's videogenerator Sora, die met kunstmatige intelligentie filmpjes kan maken uit een simpele tekstbeschrijving, heeft een bewogen week achter de rug: hoewel de technologie zelf indrukwekkend is, bleek de bijbehorende sociale app - een soort AI-versie van TikTok - niet genoeg mensen te boeien en wordt nu stopgezet. Tegelijkertijd benadrukt OpenAI dat veiligheid de fundering vormt waarop ze verder bouwen met hun nieuwste model, Sora 2. Google kwam met Lyria 3 Pro op de proppen, waarmee je langere muziekstukken kunt laten maken door AI. Aan de beeldkant lanceerde Luma Labs het model Uni-1, dat als een architect eerst nadenkt over wat je écht bedoelt voordat het een afbeelding bouwt - in plaats van blindelings pixels bij elkaar te gooien. Kort samengevat: AI wordt steeds beter in het creëren van beeld, video en geluid, maar de grote uitdaging blijft om er ook iets van te maken dat mensen écht willen gebruiken.

### Creating with Sora Safely

OpenAI Blog

To address the novel safety challenges posed by a state-of-the-art video model as well as a new social creation platform,

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

we've built Sora 2 and the Sora app with safety at the foundation. Our approach is anchored in concrete protections.

<https://openai.com/sora-2>

## Lyria 3 Pro: Create longer tracks in more Google products

Google AI

Sizzle video showing new capabilities from Lyria 3 Pro

<https://blog.google/technology/ai/lyria-3-pro>

## OpenAI's Sora was the creepiest app on your phone - now it's shutting down

TechCrunch

Though the underlying Sora 2 video- and audio-generation model is scarily impressive, there was not sustained interest in an AI-only social feed.

<https://techcrunch.com/2024/03/26/openais-sora-was-the-creepiest-app-on-your-phone-now-its-shutting-down/>

## Luma Labs Launches Uni-1: The Autoregressive Transformer Model that Reasons through Intentions Before Generating Images

MarkTechInsights

In the field of generative AI media, the industry is transitioning from purely probabilistic pixel synthesis toward models capable of structural reasoning. Luma Labs has just released Uni-1, a foundational image model designed to address the 'intent gap' inherent in standard diffusion pipelines. By implementing a reasoning phase prior to generation, Uni-1 shifts the workflow [and] ...

<https://www.marktechinsights.com/blog/luma-labs-launches-uni-1-the-autoregressive-transformer-model-that-reasons-through-intentions-before-generating-images/>

## Top 10 YouTube Channels to Learn Machine Learning

Analytics Vidya

With so much happening in AI and machine learning today, figuring out where to start can feel overwhelming. Different learners prefer different approaches! Some want visuals, others prefer coding. Some prefer short form, others lean toward long-form learning. While many simply want a clear path into ML. This article is here to fix that. Instead [and] The post Top 10 YouTube Channels to Learn M...

<https://www.analyticsvidya.com/2024/03/top-10-youtube-channels-to-learn-machine-learning/>

## Large Language Models (18 articles)

### Samenvatting

Grote techbedrijven als OpenAI, Tencent en Meta sleutelen volop aan hun taalmodellen - de digitale breinen achter chatbots zoals ChatGPT. OpenAI bouwde nieuwe vangrails in om jongeren beter te beschermen online en maakte van ChatGPT een soort persoonlijk winkelassistent die producten kan vergelijken en aanbevelen. Tegelijkertijd ontstond flinke ophef: het Amerikaanse leger wil AI inzetten voor defensie, wat leidde tot ruzie met Anthropic én het grootste straatprotest tegen AI ooit in Londen. Aan de technische kant vonden onderzoekers slimme manieren om deze modellen zuiniger en sneller te maken - een beetje zoals een auto die opeens twee keer zo ver rijdt op dezelfde tank brandstof. Tot slot werd ook duidelijk dat steeds meer bedrijven, van China tot India, meedoen aan de wereldwijde wedloop om de beste AI te bouwen.

## Helping developers build safer AI experiences for teens

OpenAI Blog

OpenAI releases prompt-based teen safety policies for developers using gpt-oss-safeguard, helping moderate age-specific

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

risks in AI systems.

<https://open>

## Powering product discovery in ChatGPT

OpenAI B

ChatGPT introduces richer, visually immersive shopping powered by the Agentic Commerce Protocol, enabling product discovery, side-by-side comparisons, and merchant integration.

<https://open>

## The AI Hype Index: AI goes to war

MIT Techn

AI is at war. Anthropic and the Pentagon feuded over how to weaponize Anthropic's AI model Claude; then OpenAI swept the Pentagon off its feet with an "opportunistic and sloppy" deal. Users quit ChatGPT in droves. People marched through London in the biggest protest against AI to date. If you're keeping score, Anthropic-the company founded&#8230;

<https://www>

## Mercor competitor Deccan AI raises \$25M, sources experts from India

TechCrun

Deccan AI concentrates its workforce in India to manage quality in a fast-growing but fragmented AI training market.

<https://tech>

## Tencent AI Open Sources Covo-Audio: A 7B Speech Language Model and Inference Pipeline for Real-Time Audio Conversations and Reasoning

MarkTech

Tencent AI Lab has released Covo-Audio, a 7B-parameter end-to-end Large Audio Language Model (LALM). The model is designed to unify speech processing and language intelligence by directly processing continuous audio inputs and generating audio outputs within a single architecture. System Architecture The Covo-Audio framework consists of four primary components designed for seamless cross-modal int...

<https://www>

## How to Build a Vision-Guided Web AI Agent with MolmoWeb-4B Using Multimodal Reasoning and Action Prediction

MarkTech

In this tutorial, we explore MolmoWeb, Ai2's open multimodal web agent that understands and interacts with websites directly from screenshots, without relying on HTML or DOM parsing. We set up the full environment in Colab, load the MolmoWeb-4B model with efficient 4-bit quantization, and build the exact prompting workflow that lets the model reason about [&#8230;] The post How to Build a Vision-G...

<https://www>

## NVIDIA AI Introduces PivotRL: A New AI Framework Achieving High Agentic Accuracy With 4x Fewer Rollout Turns Efficiently

MarkTech

Post-training Large Language Models (LLMs) for long-horizon agentic tasks-such as software engineering, web browsing, and complex tool use-presents a persistent trade-off between computational efficiency and model generalization. While Supervised Fine-Tuning (SFT) is computationally inexpensive, it frequently suffers from out-of-domain (OOD) performance degradation



# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

and struggles to generalize beyo...

<https://www>

## Paged Attention in Large Language Models LLMs

MarkTech

When running LLMs at scale, the real limitation is GPU memory rather than compute, mainly because each request requires a KV cache to store token-level data. In traditional setups, a large fixed memory block is reserved per request based on the maximum sequence length, which leads to significant unused space and limits concurrency. Paged Attention [230] The post Paged Attention in Large Langua...

<https://www>

## A Coding Implementation to Design Self-Evolving Skill Engine with OpenSpace for Skill Learning, Token Efficiency, and Collective Intelligence

MarkTech

In this tutorial, we explore OpenSpace, a self-evolving skill engine developed by HKUDS that makes AI agents smarter, more cost-efficient, and capable of learning from every task they perform. We walk through the complete lifecycle of OpenSpace: from installing and configuring an OpenAI model, to executing cold-start tasks where no prior skills exist, watching the [230] The post A Coding Imple...

<https://www>

## This AI Paper Introduces TinyLoRA, A 13-Parameter Fine-Tuning Method That Reaches 91.8 Percent GSM8K on Qwen2.5-7B

MarkTech

Researchers from FAIR at Meta, Cornell University, and Carnegie Mellon University have demonstrated that large language models (LLMs) can learn to reason using a remarkably small number of trained parameters. The research team introduces TinyLoRA, a parameterization that can scale down to a single trainable parameter under extreme sharing settings. Using this method on a [230] The post This AI...

<https://www>

## Meta AI's New Hyperagents Don't Just Solve Tasks-They Rewrite the Rules of How They Learn

MarkTech

The dream of recursive self-improvement in AI-where a system doesn't just get better at a task, but gets better at learning-has long been the 'holy grail' of the field. While theoretical models like the Gödel Machine have existed for decades, they remained largely impractical in real-world settings. That changed with the Darwin Gödel Machine (DGM), [230] The post Meta AI's New Hype...

<https://www>

## Vibe Coding a Private AI Financial Analyst with Python and Local LLMs

KDnugget

Learn to build an AI data analyst with Python: analyzes data, detects anomalies, and generates predictions using local LLMs.

<https://www>

## ChatLLM Review: Tired of Multiple AI Tools? Here's a Smarter All-in-One Alternative

KDnugget

Explore ChatLLM by Abacus AI, an all-in-one AI platform that brings together tools like ChatGPT, Claude, and Midjourney into

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

a single workflow. Learn about its features, pricing, and real-world use cases.

<https://www>

## Getting Started with Nanobot: Build Your First AI Agent

KDnugget

Learn how to set up Nanobot, connect it to WhatsApp, and power it with OpenAI GPT-5.3-Codex for a practical, always-on AI agent.

<https://www>

## 5 Practical Techniques to Detect and Mitigate LLM Hallucinations Beyond Prompt Engineering

ML Master

My friend who is a developer once asked an LLM to generate documentation for a payment API.

<https://mac>

## Beyond the Vector Store: Building the Full Data Layer for AI Applications

ML Master

If you look at the architecture diagram of almost any AI startup today, you will see a large language model (LLM) connected to a vector store.

<https://mac>

## Mistral Small 4: The One Model That Codes, Reasons, and Chats

Analytics

Artificial intelligence is rapidly evolving. New models emerge nearly every day, with each one attempting to be the best. In this sea of similar models, we see something new every now and then. One of such models is the new Mistral Small 4. It is an innovative AI model that is not only going to [8230;] The post Mistral Small 4: The One Model That Codes, Reasons, and Chats appeared first on Analy...

<https://www>

## PageIndex vs Traditional RAG: A Better Way to Build Document Chatbots

Analytics

What if the way we build AI document chatbots today is flawed? Most systems use RAG. They split documents into chunks, create embeddings, and retrieve answers using similarity search. It works in demos but often fails in real use. It misses obvious answers or picks the wrong context. Now there is a new approach called [8230;] The post PageIndex vs Traditional RAG: A Better Way to Build Document ...

<https://www>

## Research & Science (4 articles)

### Samenvatting

OpenAI zet al zijn kaarten op één paard: het bedrijf wil een volledig zelfstandige AI-onderzoeker bouwen die zonder menselijke hulp grote, ingewikkelde problemen kan oplossen. Google kwam met TurboQuant, een slimme techniek die het werkgeheugen van AI-modellen tot zes keer kleiner maakt - vergelijkbaar met het inpakken van een volle kledingkast in één handbagage, zonder ook maar één kledingstuk achter te laten. Het internet vergeleek het meteen met de fictieve compressietechnologie uit de serie \*Silicon Valley\*. Ondertussen werkt het team van AI-onderzoeker Yann LeCun aan betere "wereldmodellen": systemen die proberen de wereld te begrijpen zoals wij dat doen, maar die nu nog de neiging hebben om in

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

een soort luie stand te schieten waarbij ze alles op één hoop gooien. Kortom, de grote techbedrijven proberen AI niet alleen slimmer, maar ook zelfstandiger en zuiniger met geheugen te maken.

## OpenAI is throwing everything into building a fully automated researcher

MIT Techn

OpenAI is refocusing its research efforts and throwing its resources into a new grand challenge. The San Francisco firm has set its sights on building what it calls an AI researcher, a fully automated agent-based system that will be able to go off and tackle large, complex problems by itself. ??OpenAI says that this new&#230;

<https://www>

## Google unveils TurboQuant, a new AI memory compression algorithm - and yes, the internet is calling it 'Pied Piper'

TechCrun

Google's TurboQuant has the internet joking about Pied Piper from HBO's "Silicon Valley." The compression algorithm promises to shrink AI's "working memory" by up to 6x, but it's still just a lab experiment for now.

<https://tech>

## Google Introduces TurboQuant: A New Compression Algorithm that Reduces LLM Key-Value Cache Memory by 6x and Delivers Up to 8x Speedup, All with Zero Accuracy Loss

MarkTech

The scaling of Large Language Models (LLMs) is increasingly constrained by memory communication overhead between High-Bandwidth Memory (HBM) and SRAM. Specifically, the Key-Value (KV) cache size scales with both model dimensions and context length, creating a significant bottleneck for long-context inference. Google research team has proposed TurboQuant, a data-oblivious quantization framework des...

<https://www>

## Yann LeCun's New LeWorldModel (LeWM) Research Targets JEPAC Collapse in Pixel-Based Predictive World Modeling

MarkTech

World Models (WMs) are a central framework for developing agents that reason and plan in a compact latent space. However, training these models directly from pixel data often leads to &#216;representation collapse,&#217; where the model produces redundant embeddings to trivially satisfy prediction objectives. Current approaches attempt to prevent this by relying on complex heuristics: they [&#2...

<https://www>

## Robotics & Hardware (2 articles)

### Samenvatting

Robots en slimme machines kruipen steeds dieper in ons dagelijks leven, van de woonkamer tot de wolkenkrabber. First Lady Melania Trump pleit ervoor dat AI en robots een grote rol gaan spelen in het thuisonderwijs van Amerikaanse kinderen - stel je een geduldige digitale juf voor die 24 uur per dag klaarstaat. Ondertussen haalde het bedrijf Lucid Bots 20 miljoen dollar op omdat de vraag naar hun ramenwasdrones en schoonmaakrobots explosief groeit. Waar glazenwassers vroeger met ladder en spons gevaarlijk werk deden, nemen drones dat nu steeds vaker over. Kortom, robots worden als het ware de nieuwe handjes in zowel het klaslokaal als op de werkplek.

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

## Melania Trump wants a robot to homeschool your child

The first lady sees AI and robotics playing a prominent role in the future of American education.

TechCrunch

<https://techcrunch.com/2026/03/26/melania-trump-robot-homeschool/>

## Lucid Bots raises \$20M to keep up with demand for its window-washing drones

Lucid Bots has seen demand accelerate over the last year for its window-cleaning drones and power-washing robots.

TechCrunch

<https://techcrunch.com/2026/03/26/lucid-bots-raises-20m/>

## General AI News (9 articles)

### Samenvatting

Google onthulde Lyria 3, een nieuw model dat muziek kan maken - denk aan een digitale componist die op commando liedjes schrijft. Tegelijk wordt de wereld van AI-"agents" steeds serieuzer: dat zijn slimme digitale assistenten die niet alleen advies geven, maar ook daadwerkelijk taken uitvoeren, zoals een complete vakantie boeken en afrekenen. Onderzoekers werken hard aan manieren om die agents betrouwbaarder te maken, onder andere door ze een beter "geheugen" te geven en te begrijpen waarom ze soms de mist ingaan. Reddit trekt ondertussen een streep in het zand tegen bots door verdachte accounts te dwingen te bewijzen dat er een echt mens achter het scherm zit. Kortom, AI wordt steeds zelfstandiger en krachtiger, en daarom groeit ook de behoefte om er grip op te houden.

## Build with Lyria 3, our newest music generation model

Google Lyria teaser

Google AI

<https://blog.google/technology/ai/lyria-3/>

## A New Framework for Evaluating Voice Agents (EVA)

<https://huggingface.co/blog/ServiceNow-AI/eva>

Hugging Face

## Build a Domain-Specific Embedding Model in Under a Day

<https://huggingface.co/blog/nvidia/domain-specific-embedding-finetune>

Hugging Face

## Agentic commerce runs on truth and context

Imagine telling a digital agent, "Use my points and book a family trip to Italy. Keep it within budget, pick hotels we've liked before, and handle the details." Instead of returning a list of links, the agent assembles an itinerary and executes the purchase. That shift, from assistance to execution, is what makes agentic AI.

MIT Technology Review

<https://www.mitre.org/technology/agentic-commerce>

## Reddit takes on the bots with new 'human verification' requirements for fishy behavior

Reddit will require suspected automated accounts to verify they're human, as it ramps up efforts to curb bot-driven spam and

TechCrunch

# Weekly AI News Digest

Generated on 2026-03-26 - Articles from the last 7 days

manipulation.

<https://tech>

## 5 Useful DIY Python Functions for Error Handling

KDnugget

Debugging Python doesn't need to be complicated. These 5 DIY functions simplify error handling and improve code reliability.

<https://www>

## How to Speed Up Slow Python Code Even If You're a Beginner

KDnugget

Slow Python code is often caused by small inefficiencies. With a few beginner-friendly techniques, you can make your programs run noticeably faster. This guide will teach you how.

<https://www>

## 7 Steps to Mastering Memory in Agentic AI Systems

ML Master

Memory is one of the most overlooked parts of agentic system design.

<https://mac>

## Why Agents Fail: The Role of Seed Values and Temperature in Agentic Loops

ML Master

In the modern AI landscape, an agent loop is a cyclic, repeatable, and continuous process whereby an entity called an AI agent — with a certain degree of autonomy — works toward a goal.

<https://mac>